

Final Report Karl-Steinbuch-Stipendium 2013/14

Gesa Stupperich, Qiu Zhiqiong, David Höppner and Manuel Sassmann

October 23, 2014

Abstract

In this article we give an technical overview of our platform for automatic and manual TEI annotation of classic chinese texts from the *Siku quanshu*. We conclude with further research directions and with preliminary results reached.

1 Introduction

The *Siku quanshu* 四庫全書 (Complete library of the four branches) was one of the last major imperial editorial project of the soon ending imperial china.¹ The sixth ruler of the Qing 清 dynasty, the long reigning Qianlong 乾隆 emperor (r. 1735–1796) issued an edict, on the very first day of the new chinese year 1772, calling on his officials to collect books from all over the empire to expand the imperial book collection. In the course of the edition process under the chief editors Ji Yun 紀昀 (*zi* Xiaolan 曉嵐, 1724–1805) and Lu Xixiong 陸錫熊 (*zi* Jiannan 健男, 1734–1792) 3471 works (in 79070 *juan*) of the 13501 works sent in where selected for inclusion. For additional 6793 titles extensive bibliographic and scholarly notes where written and published under the title *Siku quanshu zongmu tiyao* 四庫全書提要 (Catalogue (with critical abstracts) of the Complete library of the four branches).

1. The following section is based on the excellent description found in Endymion Wilkinson, *Chinese History: A New History* (Cambridge (Massachusetts): Harvard University Press, 2012), 945-954. For the political background, see R. Kent Guy, *The emperor's four Treasuries: Scholars and the State in the late Ch'ien-lung Era* (Hong Kong: Hong Kong University Press, 1987)

that collected of writings of the ending imperial china under a common cover.

TEI

2 Technical Report

2.1 Deployment and Development Environment

In the original project proposal we intended to use OpenIndiana², an illumos³ based distribution, as your base deployment and development platform. While we still consider it a better platform with the availability of better debugging and analysis tools such as dtrace and mdb, and a better filesystem in the form of zfs, we did run into problems connecting various shared-ip zones over a single network interface. As we don't wanted waste too much critical time at the start period of the project, we switched to a linux based distribution and used containers (LXC) to separate our services. For the container management we initially used pre-stable release of docker.⁴ That was a slightly bumpy road, as we encountered various bugs and interface changes. With the advent of docker version 1.0 the interface has stabilized and the platform has matured. But its design is still inferior to other container implementations like the solaris zones or freebsd jails.

We provide docker container build recipes for our postgres, exist, fuseki, rabbitmq and django services. For the developer we provide a vagrant⁵ file, which installs your deployment distribution (CoreOS) into a virtual machine and then provisions our docker containers.⁶ Thus the developer only needs to run the command: `# vagrant up` to install a local development environment. Deployment to a server is similar simple by running the docker container build recipes by hand.

2. <http://openindiana.org/>

3. <http://wiki.illumos.org/display/illumos/illumos+Home>

4. <https://www.docker.com/>

5. <https://www.vagrantup.com/>

6. CoreOS is a custom minimal distribution just to run containers <https://coreos.com/> CoreOS its similar to the illumos based SmartOS.

2.2 Web Application

The web application provides the primary interface to the user. It is the main interaction point with our services. The web application connects and interacts with the various data sources and presents the results to the user. We currently use the django framework⁷ with various additional modules, the main advantage of a python based framework resides in the availability of some third-party modules, that ease the interaction with existdb and XML documents.⁸ The user is presented with an interface to browse our collection of texts by author or title. Other services are also available to the user, like our OCR or the annotation service.

We consider the git file repository of your TEI documents are the master copy of our documents. The main reason for this stipulation resides in the UIMA component which annotates new entities on demand, but for which we currently only have a file collection reader. If an update to the master copy is applied, our XML database will reload the documents that did change.

2.2.1 Text Viewer

The user can browse TEI documents currently loaded in the XML database by title or author. If he desires to view a single document, the TEI document under question will be fetched from the XML database. Before the document is present to the user it must be XTSL transformed to rewrite the TEI tags into valid HTML tags.

Texts and single chapters can be downloaded in a plain text format or as pdf files. The user has the option to undisplay various annotations he has no interest in, for example he can hide the (modern) interpunction by pressing ”.”. Users should consult the help modal page for further shortcuts.

New annotations can also be added in the viewer. The user selects the phrase he wants to mark with the mouse and then enters a keyboard shortcut to tag the selection. The marked phrase is sent to the server and after a validation enters our RDF store as a new fact. Our RDF store is partitioned into 2 named graphs. New facts supplied by users are by default stored in an extra graph, to separate them from the core verified RDF store. A table with the available keyboard commands can be found at the documentation section of the website. Annotations can be viewed and managed in the django

7. <https://www.djangoproject.com/>

8. Specifically two packages from Emory Libraries: eulxml and eulexistdb.

admin interface if they are not automatical added to the RDF store.

2.2.2 OCR

Part of our web application is a OCR service. We use tesseract⁹ in the backend for the actually recognition. As tesseract is an open source solution the accuracy does not reach the levels of commercial offerings. Our experiments with ABBYY FineReader resulted in a much better overall recognition rate, but the licensing fee is bound to the pages scanned.¹⁰ We hope the Internet Archive upgrades they version of ABBYY some day to provide better fulltext version of they PDF documents.

2.3 SPARQL

We also run a public SPARQL (SPARQL Protocol and RDF Query Language) endpoint. This endpoint can be used via a web interface or programmatically through a client library. It provides access to our manually modeled facts and also to facts extracted from automatic annotated sources. We use the functional notation of OWL (Web Ontology Language) to generate the corresponding RDF (Resource Description Framework) triples. Currently we employ a java based RDF store.¹¹ If this solution does not scale well enough we might evaluate 4store as an alternative.¹²

2.4 UIMA

UIMA (Unstructured Information Management Architecture) started as an IBM project (Watson).¹³ It was later open sourced and is now an apache project. UIMA is a architecture and a standard for annotating and enriching unstructured data. It can be used for textual sources, but its not limited to this domain. In UIMA various components interact on top of a common analysis structure (CAS) to annotate documents. At the core we find analysis engines (AE) that annotate single documents by inserting types with offsets

9. <http://code.google.com/p/tesseract-ocr/>

10. <http://www.abbyy.de/>

11. http://jena.apache.org/documentation/serving_data/index.html

12. <http://4store.org/>

13. <http://uima.apache.org/>

into the CAS. Analysis engines can share common knowledge (data structures) by using shared resource objects. This not only bundles source code into a single place but also speeds up processing as common data structures will be constructed only once and will be cached between single documents and by all engines. Before the analysis engines analyse documents, those documents must be transformed into the CAS format. This is done by the collection reader interface. Collection readers iterate through a document collection and build CAS structures for every single element. The CAS is passed then on to the analysis engines. As a final step the CAS could then be serialized again by a CAS consumer to generate a document in the original document format. In our UIMA component we follow this workflow but an additional CAS consumer extracts new information from the annotation it encounters in the newly processed document. As analysis engine can be run in parallel or sequential order, as independent or depend engines (engine B needs the annotation results from engine A) schedules are needed. We currently use a simple linear pipeline to run our analysis engines. UIMA currently provides only this functionality more complex setups needs additional scale out frameworks that are not as easy to deploy.

While UIMA provides a mature and well structured architecture for text analysis, some feature constrain the practical use.¹⁴ In our context the static type system has the greatest negative impact. As we can hold to the precondition that all the input documents to our collection reader are valid TEI documents, we would preferable create for every encountered TEI tag a corresponding UIMA type and an annotation. This would allow us to preserve all TEI tags encountered, even if dont use them yourself in your analysis engines. With a static type system as currently present in UIMA we need define all tags beforehand. As the TEI standard is quite extensive in the number of tags introduced, this is a tedious and laborious task. So we currently do not implement the whole standard but only a selection of tags our analysis engine use in their automatic annotations.

14. Thilo Götz, Jörn Kottmann, and Alexander Lang, “Quo Vadis UIMA?,” *Proceedings of the Workshop on Open Infrastructures and Analysis Frameworks for HLT* (Dublin):77–82.

2.5 Future Improvements

2.5.1 Infrastructure

Single services should be clustered to provide better availability. While CoreOS's key-value store etcd and fleet provide such capabilities, Mesos¹⁵ or Kubernetes¹⁶ likely better solutions.

2.5.2 Web Application

A full XML database is not necessary for the main functions of the site. Documents should be stored in a key value cache to provide better response times and to simplify the core functions.

2.5.3 UIMA

Besides the java version the UIMA also provides a C++ interface. A possible speed up of the interactive markup should be evaluated if we rewrite the manual annotator in C++ and trim its down to its basic functionality.

3 Date and Time Expressions

In the following section we will outline a concrete annotation example. Date time expressions are particular handy, as they play a major role in historical documents and occur in a fixed standard format.

3.1 Dynasties

In our context a dynasty is a period of time in which a succession of hereditary rulers by patrilineal descent from the dynastic founder ruled.¹⁷ The probably most famous Chinese dynasty known to the general readership is the Tang 唐 dynasty which was ruled by the Li 李 clan and lasted from 618 AD to 907 AD. An overview of all terms for imperial dynasties can be found in the literature.¹⁸ Most expressions for dynasties in chinese are polysemes. For example the aforementioned *tang* has also the meanings of 'dike' or it can

15. <http://mesos.apache.org/>

16. <https://github.com/GoogleCloudPlatform/kubernetes>

17. Wilkinson, *Chinese History*, 3.

18. Table 3 in *ibid.*, 4

function as a place name.¹⁹ This introduces the general problem of ambiguous expressions, which can only be solved by further analysis of the expression itself or/and by more context. This is a problem which solves itself for most expressions with more than two characters, as the percentage of ambiguous words drops dramatically. So the expression *Nan Tang* 南唐 and *Hou Tang* 後唐 are less ambiguous.

唐朝後唐時

If we observe the contextual occurrences of single character dynasty expression we find the following uses:

3.1.1 Dynasty + Person

A dynasty appears in prefix position before a persons name to qualify the person further.

Table 1: Dynasty expression before person name

<i>Tang Han Yu</i>	唐韓愈	Han Yu (768-824) of the Tang [dynasty]
<i>Tang Xuanzong</i>	唐玄宗	Emperor Xuan of the Tang [dynasty]
<i>Tang shushi Su Jiaqing</i>	唐術士蘇嘉慶	esoteric art expert Su Jiaqing of the Tang [dynasty]
<i>Tang xiang Pei Du</i>	唐相裴度	chancellor Pei Du (765-839) of the Tang [dynasty]

We iterate over persName annotations and check for prefix dynasty expressions.

3.1.2 Dynasty + Expression

A certain part of the particular dynasty is denoted.

Table 2: Dynasty as part of an time range

<i>Song yuqie bainian</i>	宋興且百年	the Song [dynasty] flourished since hundred years
<i>Song zhi zhongye</i>	宋之中葉	in the middle period of the Song [dynasty]
<i>zi Tang qi jin</i>	自唐迄今	from the Tang [dynasty] till today
<i>zi Tang zhi Zhou</i>	自唐至周	from the Tang [dynasty] till the Zhou [dynasty]

19. Wang Li, *Wang Li guhanyu zidian* (Beijing: Zhonghuo shuju, 2000), 117.

3.1.3 Dynasty + Era

Starting in the Han dynasty single reign periods were divided into eras. Usually the era name is followed by an expression which marks the begin, middle or end of the era.

Table 3: Dynasty with era names

<i>Tang wude chu</i>	唐武德初	at the beginning of the Tang era <i>wude</i>
<i>Tang kaiyuan zhong</i>	唐開元中	in the era <i>kaiyuan</i> of the Tang
<i>Tang kaiyuan mo</i>	唐開元末	at the end of the <i>kaiyuan</i> era of the Tang
<i>Tang kaiyuan jian</i>	唐開元間	in the <i>kaiyuan</i> era of the Tang

3.1.4 Dynasty + Era name + Year

Table 4: Dynasty with era names

<i>Tang tianbao wu nian</i>	唐天寶五年	the fifth year of <i>tianbao</i>
-----------------------------	-------	----------------------------------

3.1.5 Full Dates

If the dates are extended to the day level and behind we arrive at more complex expressions that mix different counting systems.

dates, lucky days, travel ->

eclipse of the sun

24 Solar Terms

4 General Problems

We looked at various external sources to obtain input data for our UIMA annotators. While the CBDB (China Biographical Database) Project²⁰ provides a extensive data set for persons between the 7th and 19th century, no indication of the relative importance of a person is given. Importing all 328000 records for individuals would slow down the annotation process, while inserting many possible wrong annotations. To provide a more controlled

20. <http://isites.harvard.edu/icb/icb.do?keyword=k16229>

development environment for our annotations components, we postponed the import of this data set. There are open scalability issues with our current RDF store that need to be solved before. We wrote some code to convert the SQL data provided by CBDB into OWL.

More in line with our own use of RDF as the main data representation stands the DBpedia project.²¹ The DBpedia project generates RDF triples from the data found in info boxes that occur in wikipedia articles. The generated triples can be download or queried online via a SPARQL endpoint. In general wikipedia provides are more selected set of well known individuals, but the included info boxes, which are vital for dbpedia, are distributed very uneven among the different language versions. In the chinese version most relevant for our purposes they are absend in most cases. At the current situation too much entries are without info boxes and the individual is missing in the generated triple data set.

5 Conclusion

21. <http://dbpedia.org/About>